

# BASICS OF OPTIMIZATION

Before jumping to GPU, take some baby steps



# BEFORE DIVING INTO PARALLELIZATION...



 Need to talk about some performance basics

# BEFORE DIVING INTO PARALLELIZATION...



- Need to talk about some performance basics
  - Why is tracing rays slow at all?

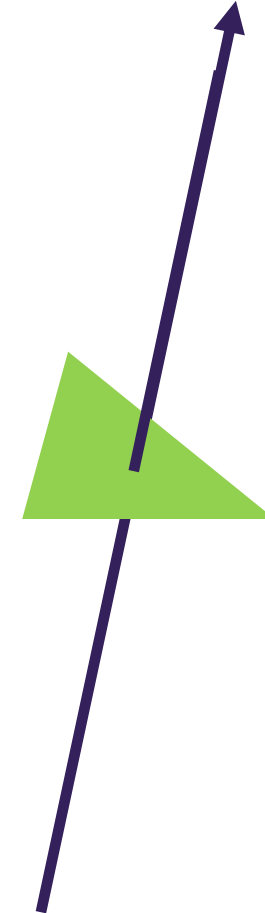
# BEFORE DIVING INTO PARALLELIZATION...

- Need to talk about some performance basics
  - Why is tracing rays slow at all?
- Consider basic ray tracing algorithm:
  - Take a ray through your scene



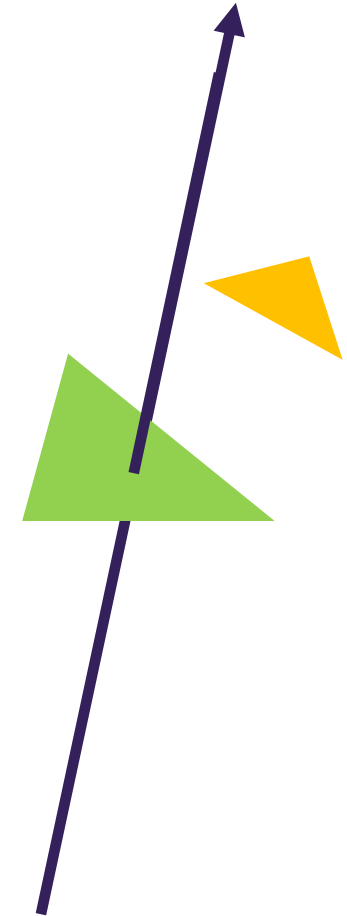
# BEFORE DIVING INTO PARALLELIZATION...

- Need to talk about some performance basics
  - Why is tracing rays slow at all?
- Consider basic ray tracing algorithm:
  - Take a ray through your scene
  - Test triangle to find intersection



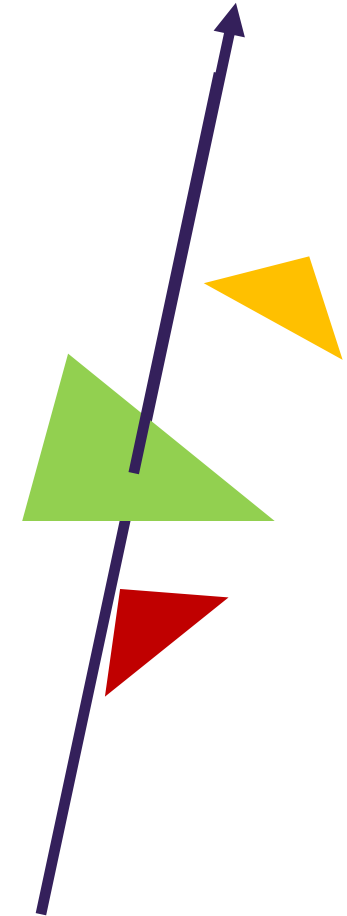
# BEFORE DIVING INTO PARALLELIZATION...

- Need to talk about some performance basics
  - Why is tracing rays slow at all?
- Consider basic ray tracing algorithm:
  - Take a ray through your scene
  - Test triangle to find intersection
    - Repeat



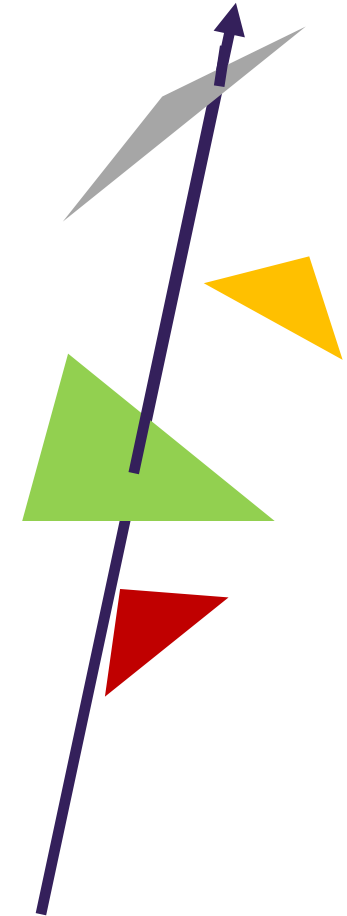
# BEFORE DIVING INTO PARALLELIZATION...

- Need to talk about some performance basics
  - Why is tracing rays slow at all?
- Consider basic ray tracing algorithm:
  - Take a ray through your scene
  - Test triangle to find intersection
    - Repeat



# BEFORE DIVING INTO PARALLELIZATION...

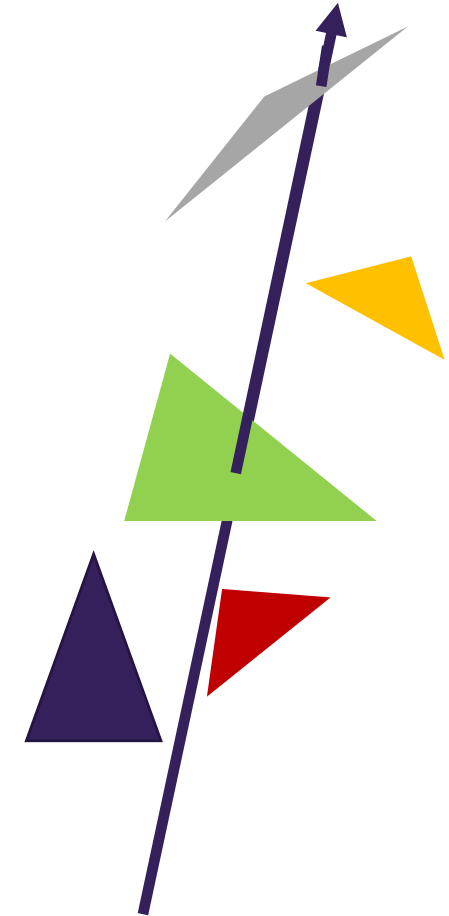
- Need to talk about some performance basics
  - Why is tracing rays slow at all?
- Consider basic ray tracing algorithm:
  - Take a ray through your scene
  - Test triangle to find intersection
    - Repeat





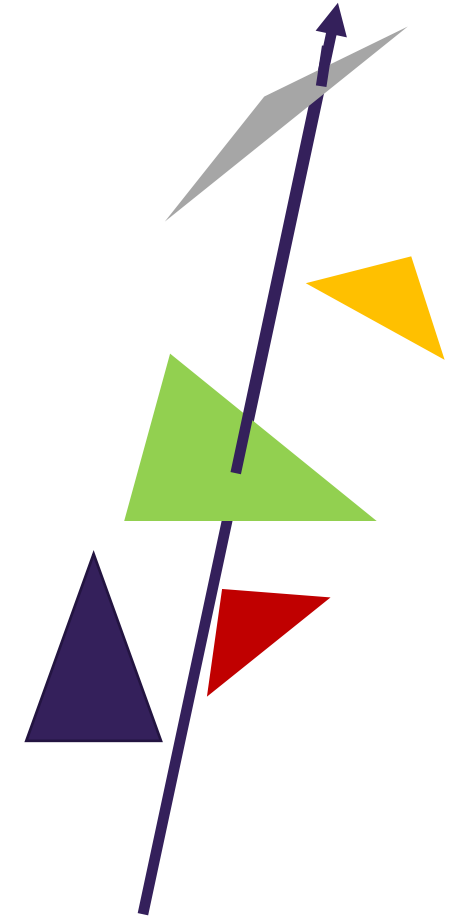
# BEFORE DIVING INTO PARALLELIZATION...

- Need to talk about some performance basics
  - Why is tracing rays slow at all?
- Consider basic ray tracing algorithm:
  - Take a ray through your scene
  - Test triangle to find intersection
    - Repeat



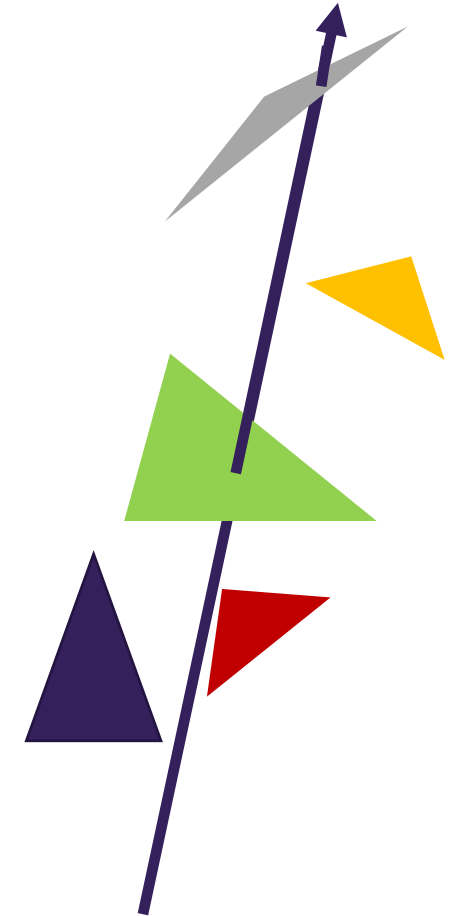
# BEFORE DIVING INTO PARALLELIZATION...

- Need to talk about some performance basics
  - Why is tracing rays slow at all?
- Consider basic ray tracing algorithm:
  - Take a ray through your scene
  - Test triangle to find intersection
    - Repeat
  - How do you know when you're done?



# BEFORE DIVING INTO PARALLELIZATION...

- Need to talk about some performance basics
  - Why is tracing rays slow at all?
- Consider basic ray tracing algorithm:
  - Take a ray through your scene
  - Test triangle to find intersection
    - Repeat
  - How do you know when you're done?
    - When you've tested every triangle?



# BEFORE DIVING INTO PARALLELIZATION...

- Need to talk about some performance basics
  - Why is tracing rays slow at all?
- Consider basic ray tracing algorithm:
  - Take a ray through your scene
  - Test triangle to find intersection
    - Repeat
  - How do you know when you're done?
    - When you've tested every triangle?
    - Very expensive...
    - Every ray could test, 1 million (or more) triangles



# WHAT'S OUR COMPUTATION BUDGET?



# WHAT'S OUR COMPUTATION BUDGET?



- Let's be easy on ourselves:
  - Target just 1920 x 1080 at 60 fps

# WHAT'S OUR COMPUTATION BUDGET?



- Let's be easy on ourselves:
  - Target just 1920 x 1080 at 60 fps
  - We need 125 million pixels per second!



# WHAT'S OUR COMPUTATION BUDGET?

- Let's be easy on ourselves:
  - Target just 1920 x 1080 at 60 fps
  - We need 125 million pixels per second!
- With a ~10 TFLOP state-of-the-art GPU
  - If tracing one ray per pixel...
  - About 80,000 flops per ray





# WHAT'S OUR COMPUTATION BUDGET?

- Let's be easy on ourselves:
  - Target just 1920 x 1080 at 60 fps
  - We need 125 million pixels per second!
- With a ~10 TFLOP state-of-the-art GPU
  - If tracing one ray per pixel...
  - About 80,000 flops per ray
- An optimized triangle intersection: ~10 flops
  - Can afford ***at most*** 8,000 intersections per ray



# WHAT'S OUR COMPUTATION BUDGET?

- Let's be easy on ourselves:
  - Target just 1920 x 1080 at 60 fps
  - We need 125 million pixels per second!
- With a ~10 TFLOP state-of-the-art GPU
  - If tracing one ray per pixel...
  - About 80,000 flops per ray
- An optimized triangle intersection: ~10 flops
  - Can afford ***at most*** 8,000 intersections per ray
- Conclusion: ***Don't test every triangle!***

# KEY PRINCIPAL TO OPTIMIZATION:

 Make the ***common*** case fast





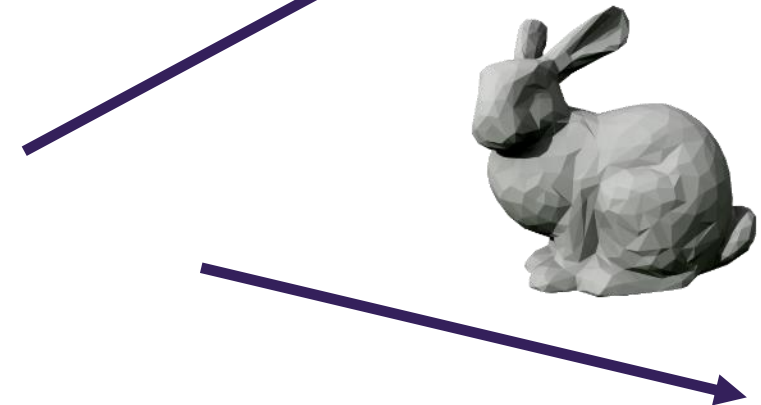
## KEY PRINCIPAL TO OPTIMIZATION:

- Make the ***common*** case fast
- Common case in ray tracing?
  - Ray does not intersect a triangle



# KEY PRINCIPAL TO OPTIMIZATION:

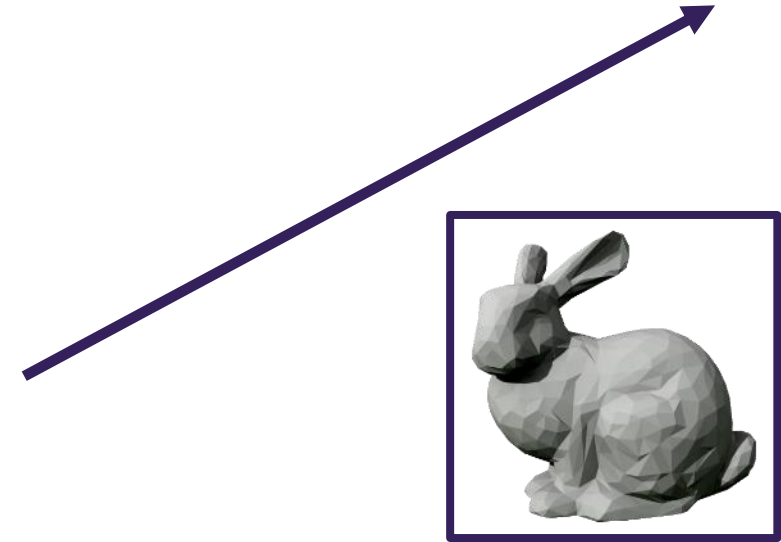
- Make the **common** case fast
- Common case in ray tracing?
  - Ray does not intersect a triangle
  - For any mesh, ray typically misses mesh





# KEY PRINCIPAL TO OPTIMIZATION:

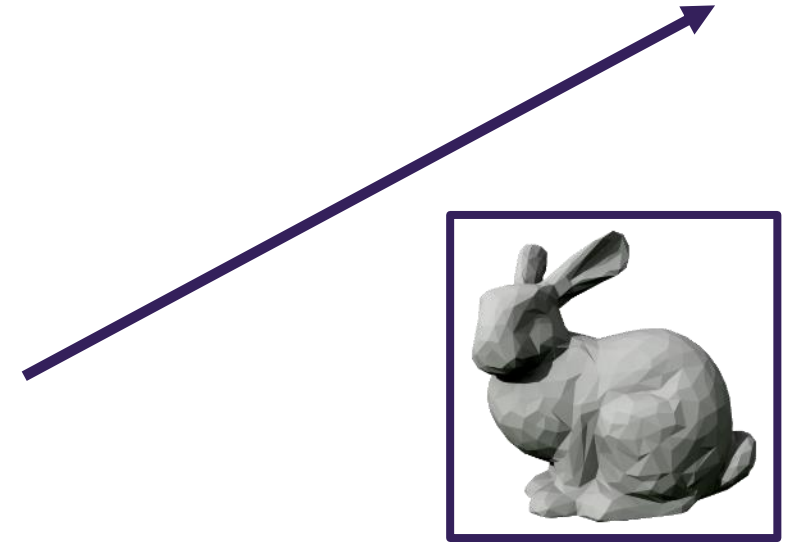
- Make the **common** case fast
- Common case in ray tracing?
  - Ray does not intersect a triangle
  - For any mesh, ray typically misses mesh
- Perhaps:
  - First intersect a mesh bounding box





# KEY PRINCIPAL TO OPTIMIZATION:

- ◆ Make the ***common*** case fast
- ◆ Common case in ray tracing?
  - Ray does not intersect a triangle
  - For any mesh, ray typically misses mesh
- ◆ Perhaps:
  - First intersect a mesh bounding box
  - Most rays avoid testing thousands of triangles
  - But, extra box test when hit bunny



# KEY PRINCIPAL TO OPTIMIZATION:

What if you have thousands of bunnies?

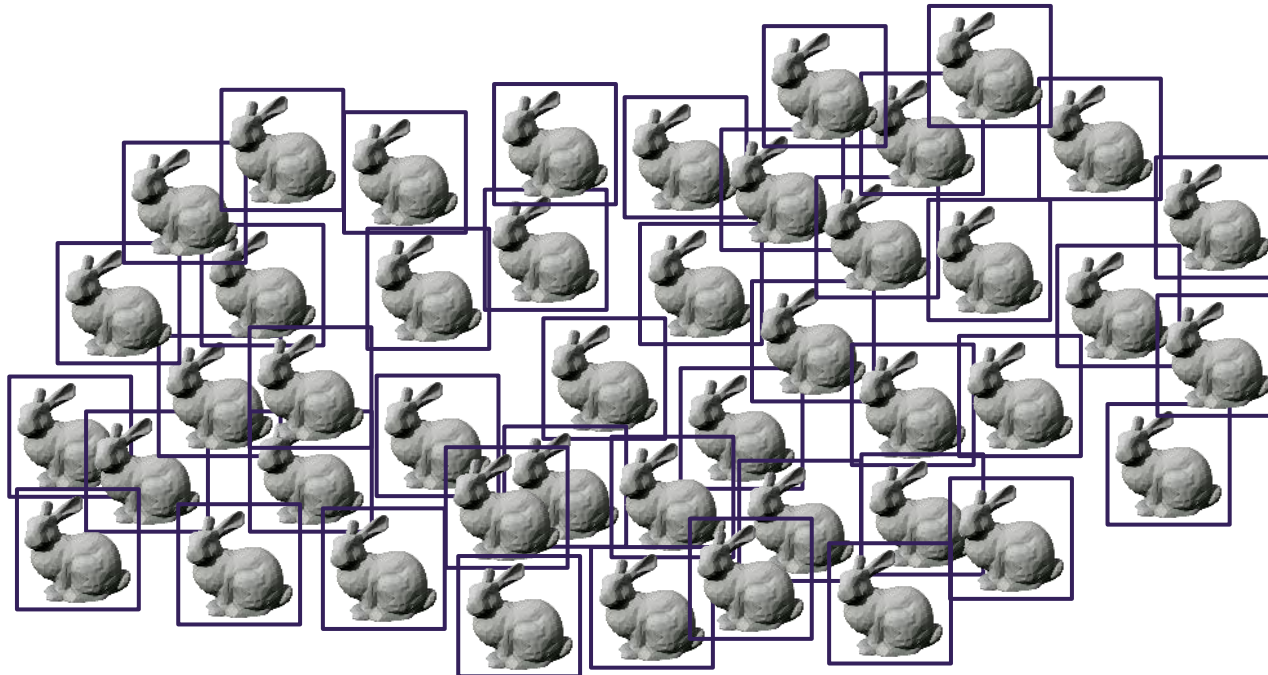






# KEY PRINCIPAL TO OPTIMIZATION:

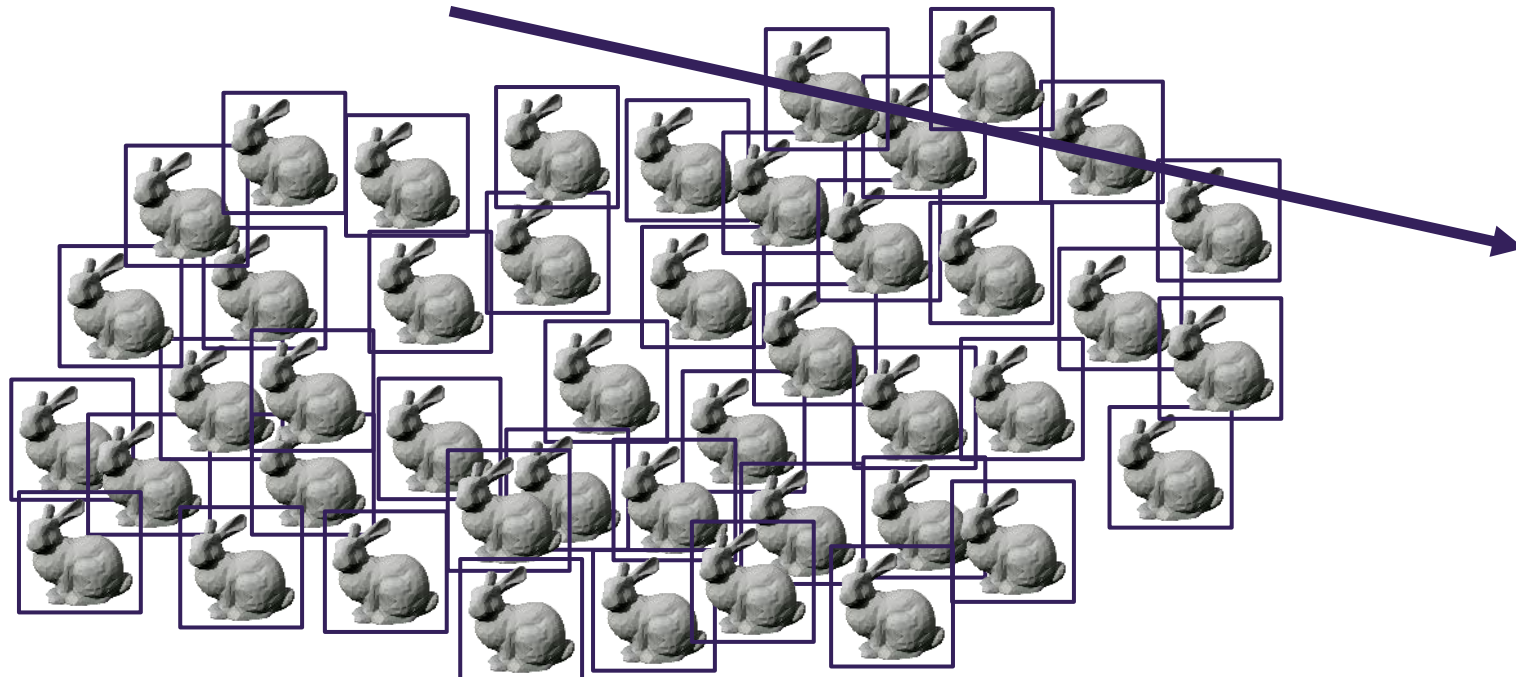
What if you have thousands of bunnies?





# KEY PRINCIPAL TO OPTIMIZATION:

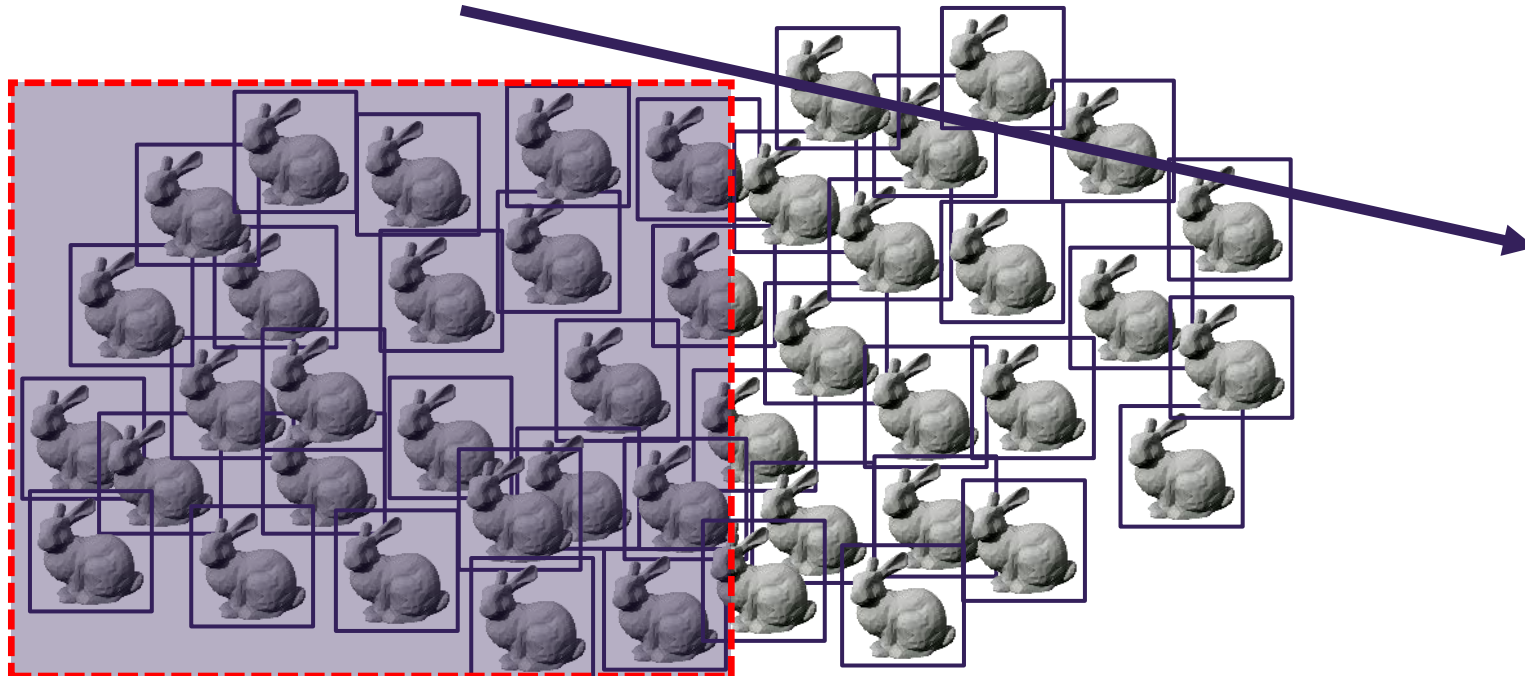
- What if you have thousands of bunnies?
  - Common case: Ray misses most bunnies





# KEY PRINCIPAL TO OPTIMIZATION:

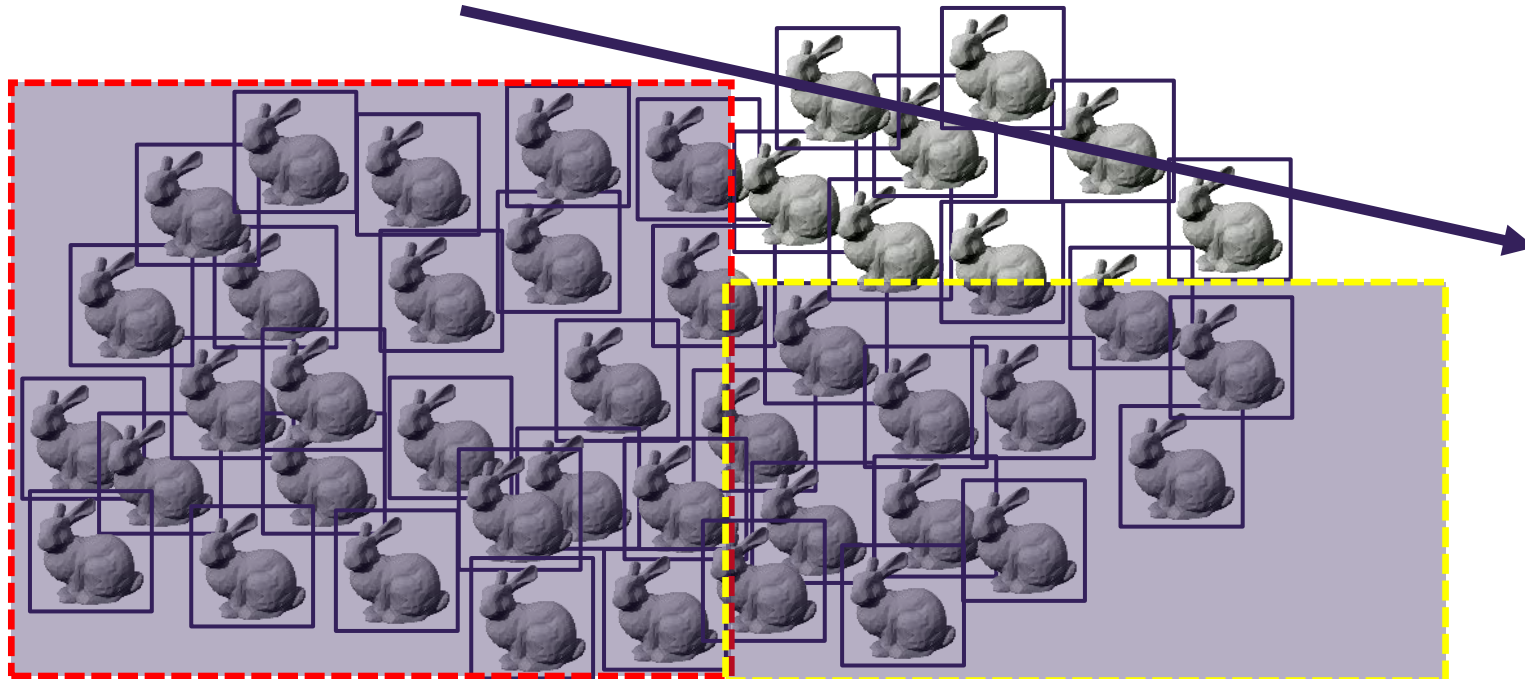
- What if you have thousands of bunnies?
  - Common case: Ray misses most bunnies
  - Can skip testing this half...





# KEY PRINCIPAL TO OPTIMIZATION:

- What if you have thousands of bunnies?
  - Common case: Ray misses most bunnies
  - Can skip testing this half... and this quarter... with a few more boxes





# KEY PRINCIPAL TO OPTIMIZATION:

- Build a tree of bounding boxes
  - Known as a “bounding volume hierarchy” or BVH



## KEY PRINCIPAL TO OPTIMIZATION:

- Build a tree of bounding boxes
  - Known as a “bounding volume hierarchy” or BVH
- When using a principled tree build
  - Reduces number of required intersections
  - From  $O(N)$  to  $O(\log N)$



# KEY PRINCIPAL TO OPTIMIZATION:

- ◆ Build a tree of bounding boxes
  - Known as a “bounding volume hierarchy” or BVH
- ◆ When using a principled tree build
  - Reduces number of required intersections
  - From  $O(N)$  to  $O(\log N)$
- ◆ With a binary tree, 1 million ray-triangle tests becomes:
  - Around 20 ray-box tests
  - A few ray-triangle tests in leaf nodes

# KEY PRINCIPAL TO OPTIMIZATION:



Production ray tracers ***always*** use some acceleration structure





## KEY PRINCIPAL TO OPTIMIZATION:

- Production ray tracers ***always*** use some acceleration structure
- But, ***which*** structure? How do you best build it?
  - Literally decades of research



# KEY PRINCIPAL TO OPTIMIZATION:

- Production ray tracers ***always*** use some acceleration structure
- But, ***which*** structure? How do you best build it?
  - Literally decades of research
  - Continuing to today (e.g., “Wide BVH Traversal with a Short Stack,” Vaidyanathan et al. 2019)



# KEY PRINCIPAL TO OPTIMIZATION:

- ◆ Production ray tracers ***always*** use some acceleration structure
- ◆ But, ***which*** structure? How do you best build it?
  - Literally decades of research
  - Continuing to today (e.g., “Wide BVH Traversal with a Short Stack,” Vaidyanathan et al. 2019)
- ◆ When starting real-time ray tracing, best bet:
  - Use someone else’s code
  - Quality of your BVH easily affects performance by 2x, 3x, or >10x
    - Varies per scene!
  - Luckily most APIs will build structure

